

Sampling_from_a_Biased_Population

October 18, 2019

0.1 Sampling from a Biased Population

In this tutorial we will go over some code that recreates the visualizations in the Interactive Sampling Distribution Demo. This demo looks at a hypothetical problem that illustrates what happens when we sample from a biased population and not the entire population we are interested in. This tutorial assumes that you have seen that demo, for context, and understand the statistics behind the graphs.

```
In [ ]: # Import the packages that we will be using for the tutorial
import numpy as np # for sampling for the distributions
import matplotlib.pyplot as plt # for basic plotting
import seaborn as sns; sns.set() # for plotting of the histograms

# Recreate the simulations from the video
mean_uofm = 155
sd_uofm = 5
mean_gym = 185
sd_gym = 5
gymperc = .3
totalPopSize = 40000

# Create the two subgroups
uofm_students = np.random.normal(mean_uofm, sd_uofm, int(totalPopSize * (1 - gymperc)))
students_at_gym = np.random.normal(mean_gym, sd_gym, int(totalPopSize * (gymperc)))

# Create the population from the subgroups
population = np.append(uofm_students, students_at_gym)

# Set up the figure for plotting
plt.figure(figsize=(10,12))

# Plot the UofM students only
plt.subplot(3,1,1)
sns.distplot(uofm_students)
plt.title("UofM Students Only")
plt.xlim([140,200])

# Plot the Gym Goers only
```

```

plt.subplot(3,1,2)
sns.distplot(students_at_gym)
plt.title("Gym Goers Only")
plt.xlim([140,200])

# Plot both groups together
plt.subplot(3,1,3)
sns.distplot(population)
plt.title("Full Population of UofM Students")
plt.axvline(x = np.mean(population))
plt.xlim([140,200])

plt.show()

```

0.2 What Happens if We Sample from the Entire Population?

We will sample randomly from all students at the University of Michigan.

```

In [ ]: # Simulation parameters
numberSamps = 5000
sampSize = 50

# Get the sampling distribution of the mean from only the gym
mean_distribution = np.empty(numberSamps)
for i in xrange(numberSamps):
    random_students = np.random.choice(population, sampSize)
    mean_distribution[i] = np.mean(random_students)

# Plot the population and the biased sampling distribution
plt.figure(figsize = (10,8))

# Plotting the population again
plt.subplot(2,1,1)
sns.distplot(population)
plt.title("Full Population of UofM Students")
plt.axvline(x = np.mean(population))
plt.xlim([140,200])

# Plotting the sampling distribution
plt.subplot(2,1,2)
sns.distplot(mean_distribution)
plt.title("Sampling Distribution of the Mean Weight of Gym Goers")
plt.axvline(x = np.mean(population))
plt.axvline(x = np.mean(mean_distribution), color = "black")
plt.xlim([140,200])

plt.show()

```

0.3 What Happens if We take a Non-Representative Sample?

What happens if I only go to the gym to get the weight of individuals, and I don't sample randomly from all students at the University of Michigan?

```
In [ ]: # Simulation parameters
        numberSamps = 5000
        sampSize = 3

        # Get the sampling distribution of the mean from only the gym
        mean_distribution = np.empty(numberSamps)
        for i in xrange(numberSamps):
            random_students = np.random.choice(students_at_gym, sampSize)
            mean_distribution[i] = np.mean(random_students)

        # Plot the population and the biased sampling distribution
        plt.figure(figsize = (10,8))

        # Plotting the population again
        plt.subplot(2,1,1)
        sns.distplot(population)
        plt.title("Full Population of UofM Students")
        plt.axvline(x = np.mean(population))
        plt.xlim([140,200])

        # Plotting the sampling distribution
        plt.subplot(2,1,2)
        sns.distplot(mean_distribution)
        plt.title("Sampling Distribution of the Mean Weight of Gym Goers")
        plt.axvline(x = np.mean(population))
        plt.axvline(x = np.mean(students_at_gym), color = "black")
        plt.xlim([140,200])

        plt.show()
```